

DEVOPS WROCŁAW MEETUP #7

Prometheus 101: get on the hype train

Michał Bolek - OPS@Cogniance



Cogniance
Your Intelligence. **Multiplied.**



TODO:

- ŻŻP
- What is prometheus?
- Architecture
- Short comparison
- Lifelike topics
- Demo



PROMETHEUS: YESTERDAY'S TECHNOLOGY TODAY



Even though Borgmon remains internal to Google, the idea of treating time-series data as a data source for generating alerts is now accessible to everyone through those open source tools like Prometheus [...]



WHAT IS PROMETHEUS ?

- www.prometheus.io
- 'quite' old monitoring and alerting toolkit - started around 2012
- built at SoundCloud (by some of the exgooglers)
- scrapes and stores time series data
- and then uses them for monitoring
- pull based (but can push if needed)
- standalone
- PromQL
- joined the Cloud Native Computing Foundation in 2016 just after Kubernetes...




TRENDS IN SEARCH:

Interest over time

Google Trends

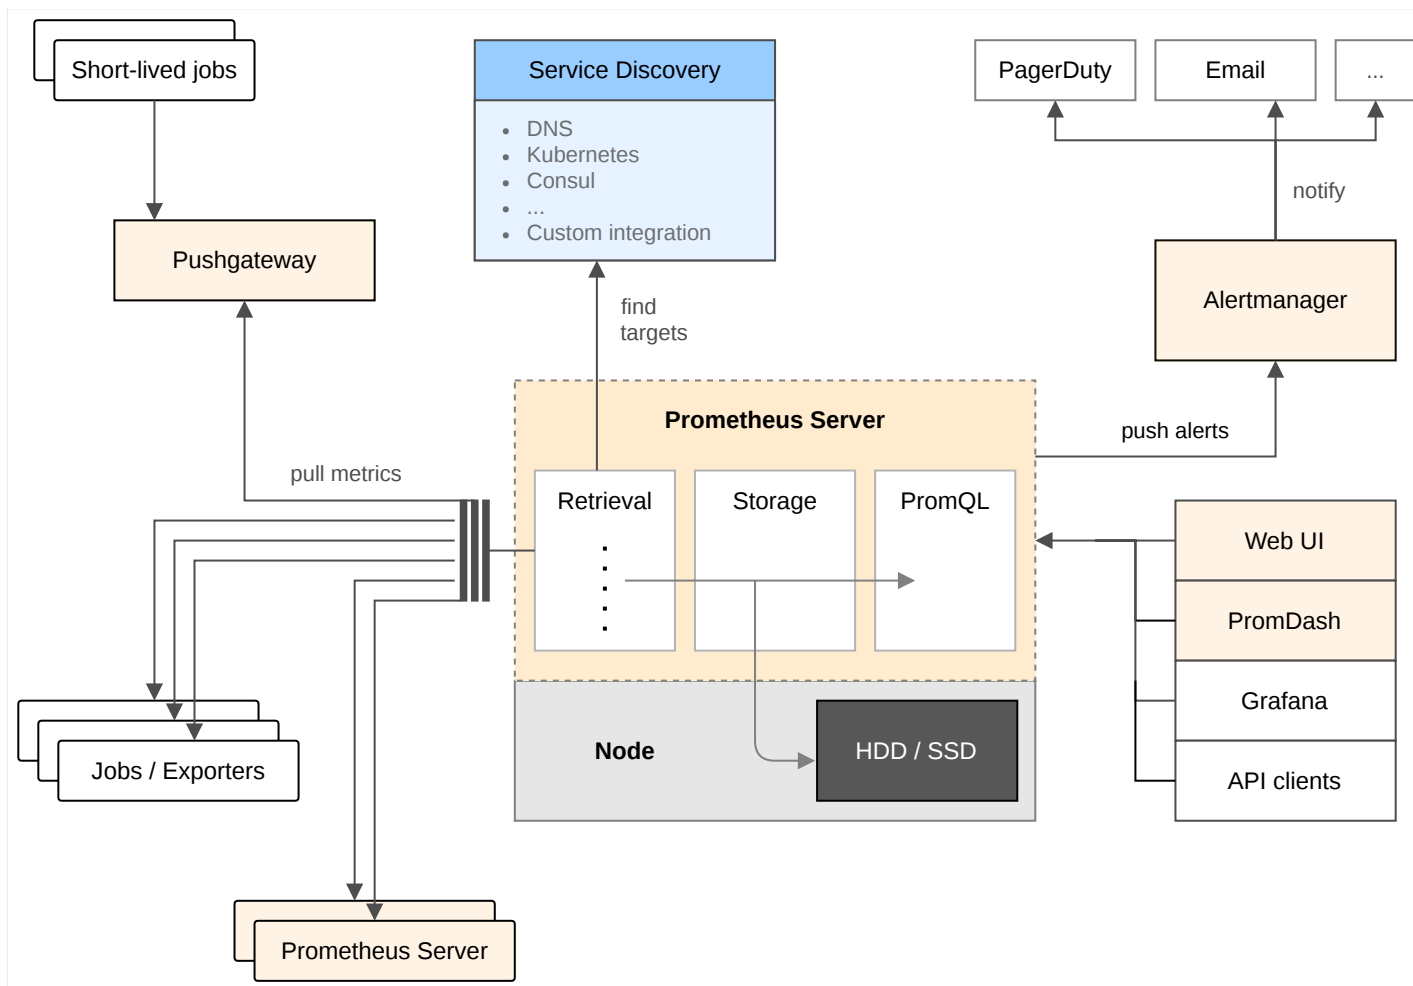
● prometheus ● nagios



Worldwide. Past 5 years.



ARCHITECTURE



Structure of timeseries

	:	:	:	:	:	:	:	:	:	:
	0	0	0	0	0	0	0	0	0	
:	0	0	0	0	0	0	0	0	0	
now - 2 Δt	0	0	0	0	0	0	0	0	0	
now - Δt	0	0	0	0	0	0	0	0	0	
now	0	0	0	0	0	0	0	0	0	...
	label1	label2	label3	label4	...					

api_http_requests_total{method="GET",
endpoint="/api/", status="200"} - 123



PROMETHEUS SERVER

- scrapes endpoints' /metrics
- saves data
- run queries



SERVICE DISCOVERY

- DNS
- consul
- kubernetes
- etcd
- EC2
- marathon
- static file



BASIC PROMETHEUS CONFIG

```
global:
  scrape_interval:      15s # By default, scrape targets every 15s.

  # Attach these labels to any time series or alerts when communicating
  # with external systems (federation, remote storage, Alertmanager).
  external_labels:
    monitor: 'codelab-monitor'

# A scrape configuration containing exactly one endpoint to scrape:
# Here it's Prometheus itself.
scrape_configs:
  # The job name is added as a label `job=`
  - job_name: 'prometheus'

# Override and scrape targets from this job every 5 seconds.
  scrape_interval: 5s

  static_configs:
    - targets: ['localhost:9090']
```



ALERTMANAGER



CONFIG

```
global:

route:
  group_by: ['alertname', 'cluster']
  group_wait: 30s
  group_interval: 5m
  repeat_interval: 1h
  continue: true
  receiver: 'default-slack'

receivers:
- name: 'default-slack'
  slack_configs:
    - api_url: 'https://hooks.slack.com/services/wc4l3n13pR4wD21wy70k3n'
      send_resolved: true
      channel: '#monitoring'
      username: 'Prometheus'
      icon_emoji: ':crystal_ball:'
      text: {% raw %}"{{ .CommonAnnotations.description }}" {% endraw %}
```





Prometheus APP 11:27 AM

[FIRING:1] BackendNotificationFail (https://
- notifications issue



Prometheus APP 11:49 AM

[FIRING:5] InstanceDown (maven)



Prometheus APP 11:58 AM

[FIRING:3] InstanceLowMem (node maven)



Prometheus APP 12:32 PM

[FIRING:1] BackendNotificationFail (http
- notifications issue



Prometheus APP 12:54 PM

[FIRING:5] InstanceDown (maven)



Prometheus APP 1:03 PM

[FIRING:3] InstanceLowMem (node maven)



Prometheus APP 1:26 PM

[RESOLVED] InstanceSlowIO (10.2.1.188:9100 maven)
10.2.1.188:9100 slow IO - 1234.4

[FIRING:1] InstanceSlowIO (10.2.1.188:9100 maven)
10.2.1.188:9100 slow IO - 426.13333333333333

[RESOLVED] InstanceSlowIO (10.2.1.188:9100 maven)
10.2.1.188:9100 slow IO - 426.13333333333333

[FIRING:1] BackendNotificationFail (https://
- notifications issue



EXPORTERS

- node_exporter
- Apache/Nginx/Varnish/haproxy_exporter
- MySQL/PostgreSQL/MongoDB/Redis_exporter
- AWS/Cloudflare/DockerHub/GitHub/Openweathermap
- collectd/graphite/munin/statsd_exporter
- jenkins/rtorrent/minecraft_exporter/blackbox
- JSON_exporter
- +native apps



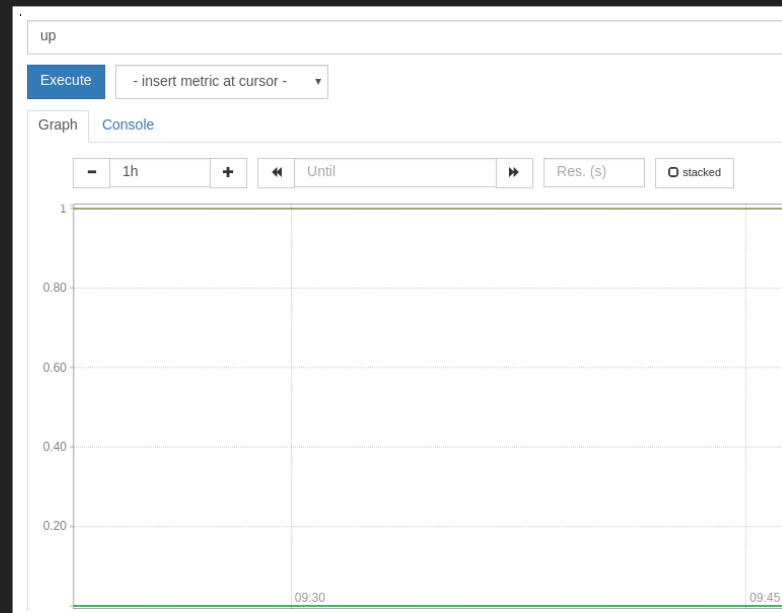
```
- job_name: 'blackbox'
  metrics_path: /probe
  params:
    module: [module_name] # Look for a HTTP 200 response.
  static_configs:
    - targets:
      - ['127.0.0.1:9999']

  relabel_configs:
    - source_labels: [__address__]
      regex: (. *?)(:80)?
      target_label: __param_target
      replacement: https://${1}
    - source_labels: [__param_target]
      target_label: instance
    - target_label: __address__
      replacement: 127.0.0.1:9115 # Blackbox exporter.
```



PROMETHEUS EXPRESSION BROWSER

up		Load time: 13ms Resolution: 14s
Execute	- insert metric at cursor - ▾	
Graph	Console	
Element	Value	
up{instance="10.2.1.184:9101",job="haproxy"}	1	
up{instance="10.2.1.151:9100",job="node"}	1	
up{instance="10.2.1.175:9100",job="node"}	1	
up{instance="10.2.1.180:9121",job="redis"}	1	
up{instance="10.2.1.185:9101",job="haproxy"}	1	
up{instance="10.2.1.180:9100",job="node"}	1	



CONSOLES

Prometheus	Alerts	PagerDuty
Overview		
HAProxy		
Blackbox		
Node		
Prometheus		

Node

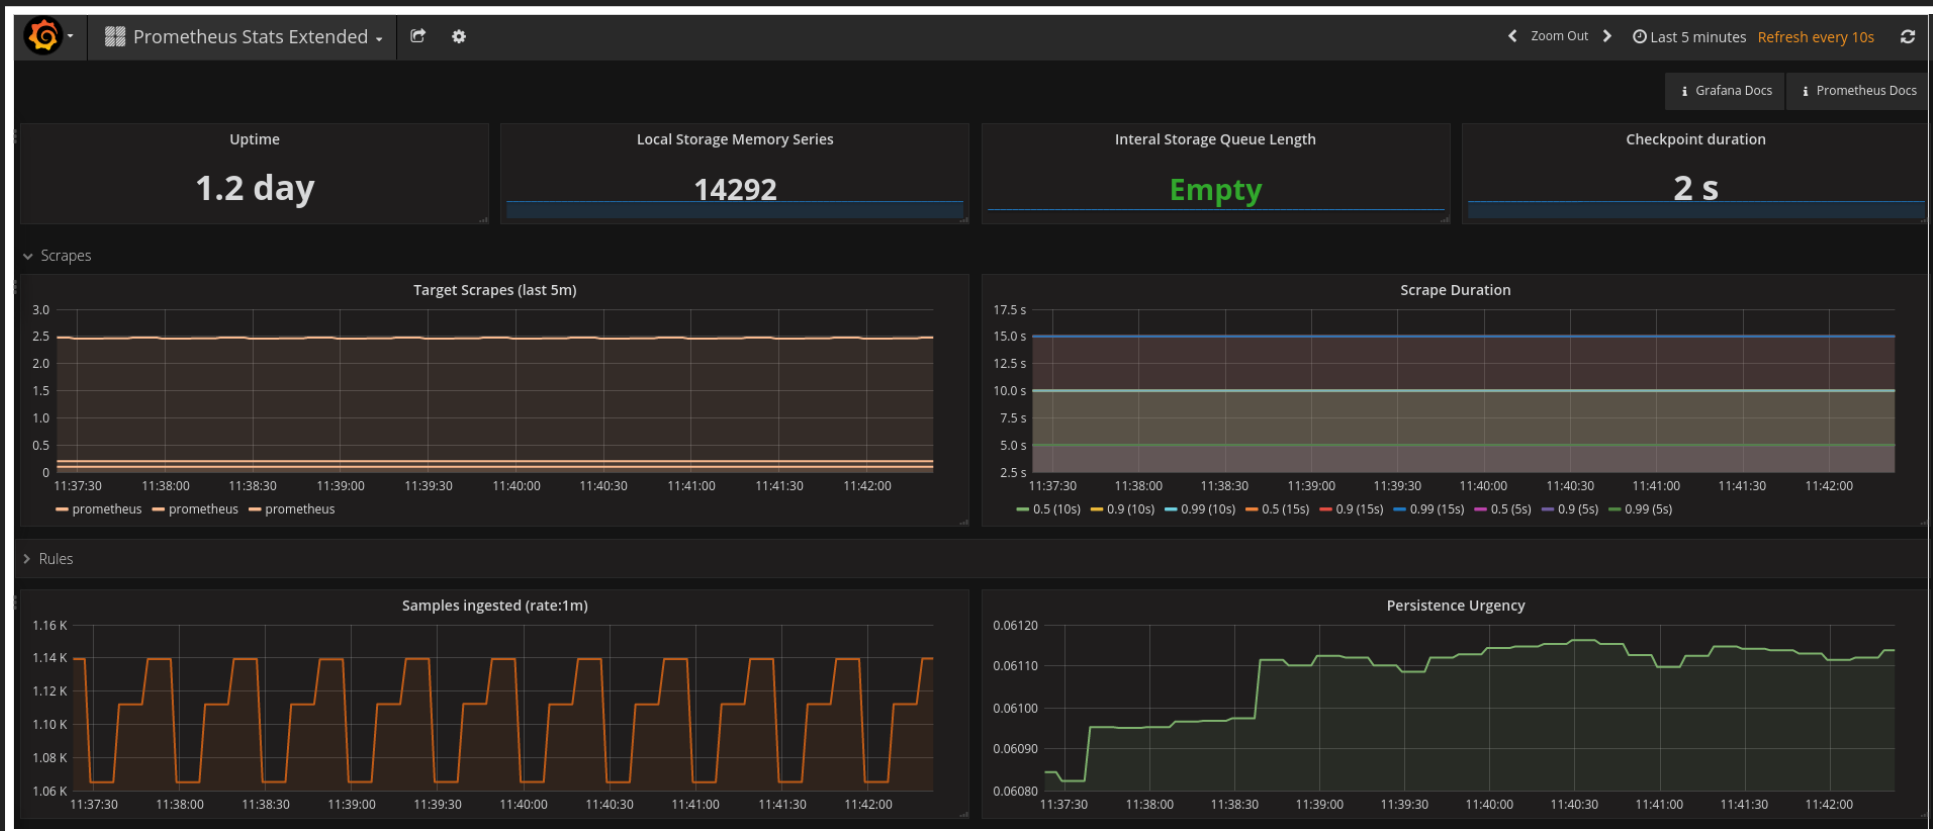
Node	Up	CPU Used	Memory Available
10.2.1.105:9100	No	-	-
10.2.1.119:9100	Yes	4.9%	2.403GiB
10.2.1.151:9100	Yes	0.1%	7.487GiB
10.2.1.174:9100	Yes	0.3%	5.528GiB
10.2.1.175:9100	Yes	1.0%	1.725GiB
10.2.1.176:9100	Yes	1.3%	870.5MiB
10.2.1.177:9100	Yes	0.3%	978.3MiB
10.2.1.178:9100	Yes	1.1%	1.726GiB
10.2.1.179:9100	Yes	1.2%	1.746GiB
10.2.1.180:9100	Yes	1.1%	1.749GiB
10.2.1.181:9100	Yes	3.6%	1.257GiB
10.2.1.182:9100	Yes	5.2%	1.268GiB
10.2.1.183:9100	Yes	4.5%	1.271GiB
10.2.1.184:9100	Yes	0.7%	1.717GiB
10.2.1.185:9100	Yes	1.2%	1.745GiB
10.2.1.188:9100	Yes	93.4%	2.688GiB



PROMDASH



GRAFANA



COMPARISON WITH GRAPHITE

Prometheus

Graphite

full monitoring

time series DB

rich metadata model

simple metadata model

LevelDB

whisper



COMPARISON WITH INFLUXDB

Prometheus

InfluxDB

full monitoring

time series DB

rich metadata model

even richer metadata model

LevelDB

distributed cluster storage



COMPARISON WITH OPENTSDDB

Prometheus

OpenTSDB

full monitoring

time series DB

rich metadata model

rich metadata model

LevelDB

Hadoop/HBase



COMPARISON WITH NAGIOS

Prometheus

Nagios

full monitoring solution

full monitoring solution

easy automation

hard automation

multi-component

monolithic



COMPARISON WITH SENSU

Prometheus

Sensu

full monitoring solution

full monitoring solution

easy automation

easy automation

multi-component


distributed



LIFELIKE TOPICS

- Pull model, is it any good?
- Nagging disks
- Expiring certs
- Need for status





Pull is ok!:)



NAGGING DISKS

Service State Information

Current Status:	WARNING (for 1d 6h 48m 55s)
Status Information:	DISK WARNING - free space: / 21505 MB (19% inode=97%):
Performance Data:	/=86231MB;90207;101483;0;112759
Current Attempt:	3/3 (HARD state)
Last Check Time:	02-23-2017 15:01:26
Check Type:	ACTIVE
Check Latency / Duration:	0.000 / 0.137 seconds
Next Scheduled Check:	02-23-2017 15:06:26
Last State Change:	02-22-2017 08:16:34
Last Notification:	02-22-2017 08:16:34 (notification 3)
Is This Service Flapping?	NO (0.00% state change)
In Scheduled Downtime?	NO
Last Update:	02-23-2017 15:05:28 (0d 0h 0m 1s ago)



NOTIFY IN ADVANCE!

```
ALERT DiskWillFillIn24Hours
  IF predict_linear(node_filesystem_free{job='node'}[1d], 24*3600) < 0
  FOR 5m
  ANNOTATIONS {
    summary = "Instance {{$labels.instance}} disk fillup in 24h!",
    description = "{{$labels.instance}}s on of the disks fill fill up in 24h!"
  }
```

[FIRING:2] DiskWillFillIn24Hours (/dev/mapper/g0-v0 xfs 10.2.1.177:9100 node maven)

10.2.1.177:9100s on of the disks fill fill up in 24h!



EXPIRING CERTS



Secure Connection Failed

i.broke.the.internet.and.all.i.got.was.this.t-shirt.phreedom.org uses an invalid security certificate.

The certificate is not trusted because the issuer certificate has expired.
The certificate expired on 9/1/2004 6:00 PM.

(Error code: sec_error_expired_issuer_certificate)

- This could be a problem with the server's configuration, or it could be someone trying to impersonate the server.
- If you have connected to this server successfully in the past, the error may be temporary, and you can try again later.

[Or you can add an exception...](#)



NEVA' AGAIN!

```
ALERT SSLCertExpiringSoon
IF probe_ssl_earliest_cert_expiry{job="blackbox"} - time() < 86400 * 30
FOR 10m
ANNOTATIONS {
    summary = "SSL Cert on {{ $labels.instance}} will expire in less than 30 days",
    description = "SSL Cert on {{ $labels.instance}} will expire in less than 30 days"
}
```



NEED FOR STATUS

```
{"smsSuccesses":[],"smsFailures":[{"version":1,"phoneNumber":"+1555123XXXX","dateTime":"2017-02-23T13:10:00Z"}, {"version":1,"phoneNumber":"+1555123XXXX","dateTime":"2017-02-23T13:13:00Z"}],"successCount":0,"failureCount":2}
```



NEED FOR STATUS

```
mvn_backend_notifications:  
  prober: http  
  timeout: 5s  
  http:  
    method: GET  
    valid_status_codes: [200]  
    no_follow_redirects: false  
    fail_if_ssl: false  
    fail_if_not_ssl: true  
    fail_if_not_matches_regexp:  
      - "\"failureCount\"\":0
```



SIMPLE EXPORTER

```
# -*- coding: UTF-8 -*-
from prometheus_client import start_http_server, Metric, REGISTRY
import json
import requests
import sys
import time

class JsonCollector(object):
    def __init__(self):
        pass

    def collect(self):
        token = 'asdadsadadasdadsasd' #apply for a token here: http://aqicn.org/data-platform/token/
        url = 'http://api.waqi.info/feed/'
        city = 'Wrocław' # Cities: http://aqicn.org/city/all/
        response = json.loads(requests.get(url+city+"/?token="+token).content.decode('Utf-8'))

        metric = Metric('aqi_airquality', 'Air Quality Index', 'gauge')
        metric.add_sample('aqi_airquality', value=response['data']['aqi'], labels={})
        yield metric

if __name__ == '__main__':
    start_http_server(int(sys.argv[1]))
    REGISTRY.register(JsonCollector())

    while True:
        time.sleep(1)
```





Wroclaw's Air Quality



< Zoom Out > Last 1 hour Refresh every 30m

Air Quality Index

312

Air Temperature

-2.02 °C


Humidity Level

88%

Pressure

1037 hPa


CO gases

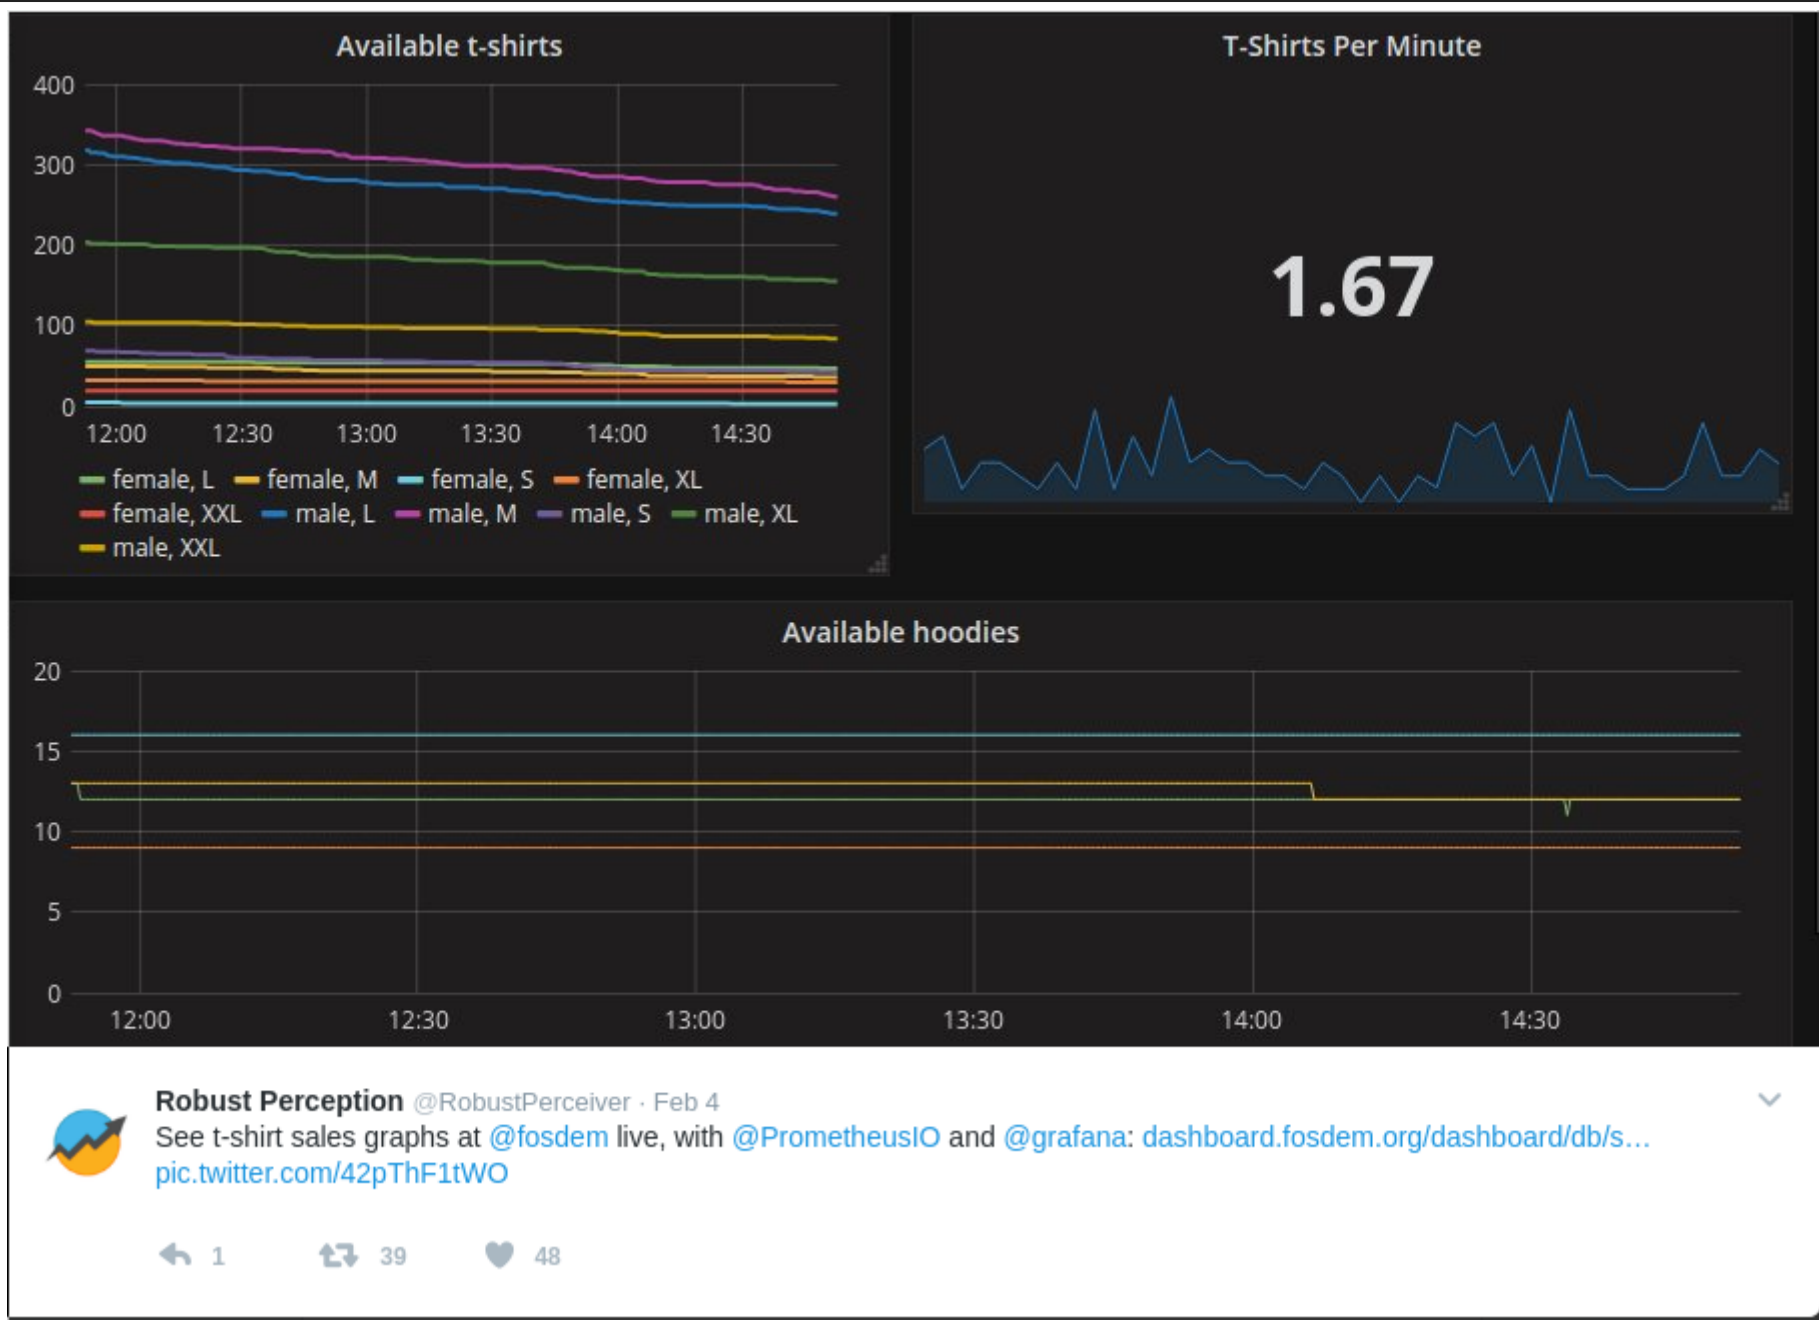


NO2 gasses



PM25





HA

ha ha ha...



brian-brazil commented on Mar 30, 2016

Member



so you could query any server, and get forwarded (http 301 or just proxy request) to a server who has the data.

This is not a feature we plan to add. Prometheus is intended for your critical monitoring, and having a full-mesh of communication between servers opens risks around cascading failures and all the complexities of having to synchronise metadata across servers.

In practice you'll almost always know which Prometheus server has the data you're looking for, so it's not a major problem in the real world.

<https://www.robustperception.io/scaling-and-federating-prometheus/>

<https://www.robustperception.io/federation-what-is-it-good-for/>



DEMO



Q's

and possibly A's



Akos Veres @puck · Feb 21

So you come in to the office, take a look at the [@PrometheusIO](#) data through [@grafana](#), spot the problem, tell the devs, apocalypse avoided.



LINKS:

- www.prometheus.io
- www.robustperception.io
- 1st post on prometheus ujeb.se/1stpost
- SRE - <https://landing.google.com/sre/book/index.html>
- Cool presentation on Prom - <http://ujeb.se/prometheus>
- Scaling to a million machines <https://www.youtube.com/watch?v=likpVWB5Lvo>
- <https://www.robustperception.io/scaling-and-federating-prometheus/>
- Prometheus-as-a-service <https://github.com/weaveworks/cortex>
- Design and Philosophy <https://www.youtube.com/watch?v=QgJbxCWRZ1s>
- Push vs Pull <http://www.boxever.com/push-vs-pull-for-monitoring/>



THANK YOU!

