



Kubernetes as a Service for everyone

Michał Jura
Linux Cloud/HA Developer
mjura@suse.com

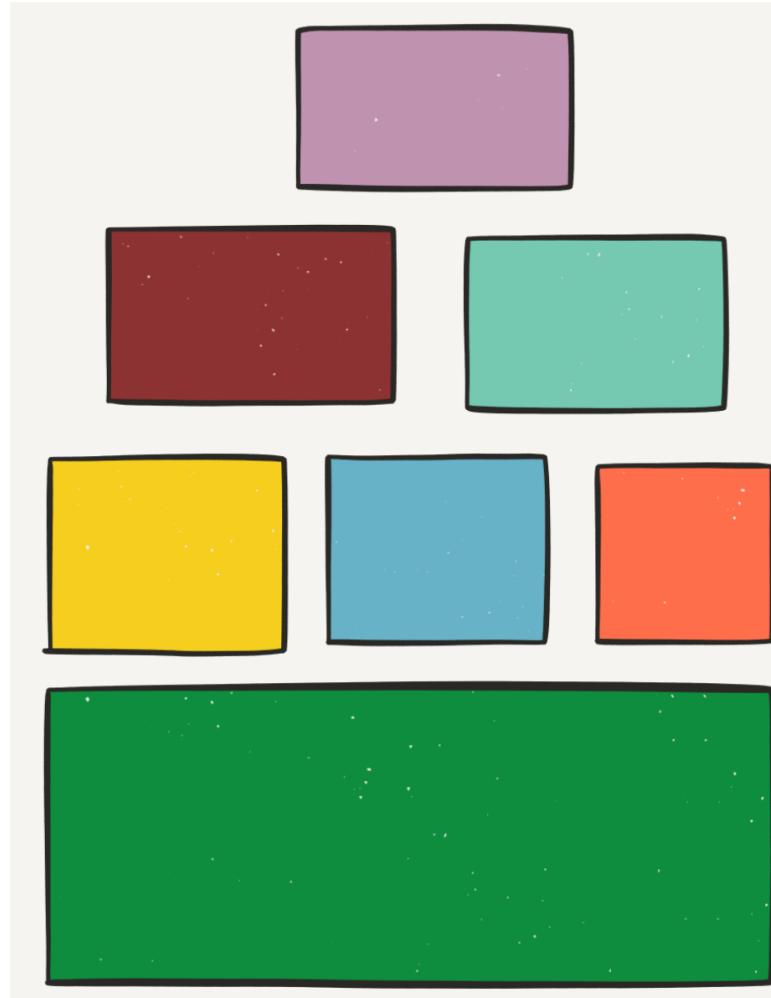
New challenges

- New age of internet and mobile applications.
- Application super portability.
- Higher complexity of cloud environments.
- Different cloud providers.
- Adoption of micro services architectures.
- Stay agile.

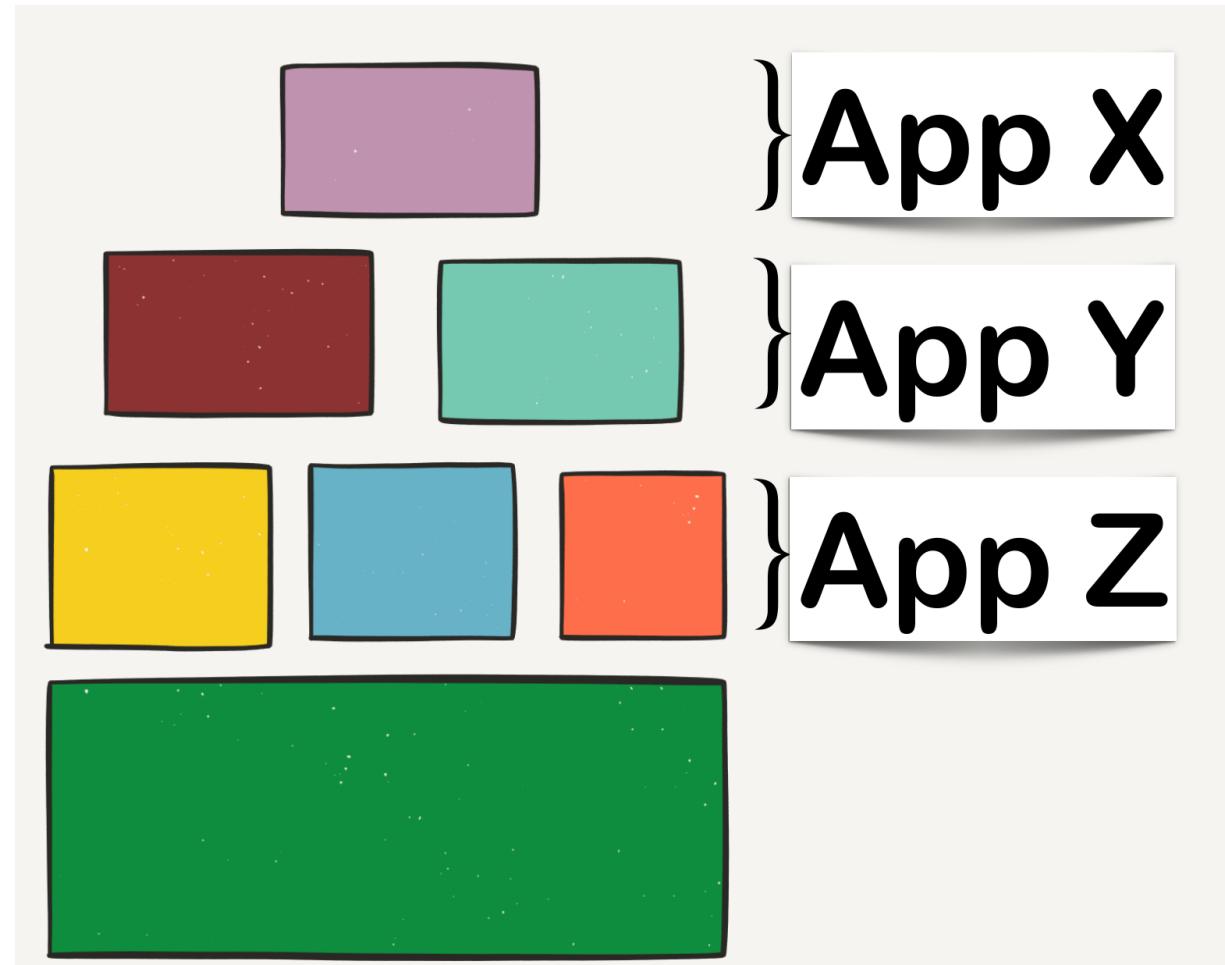
Manage applications,
not machines

**Let's talk about
application containers**

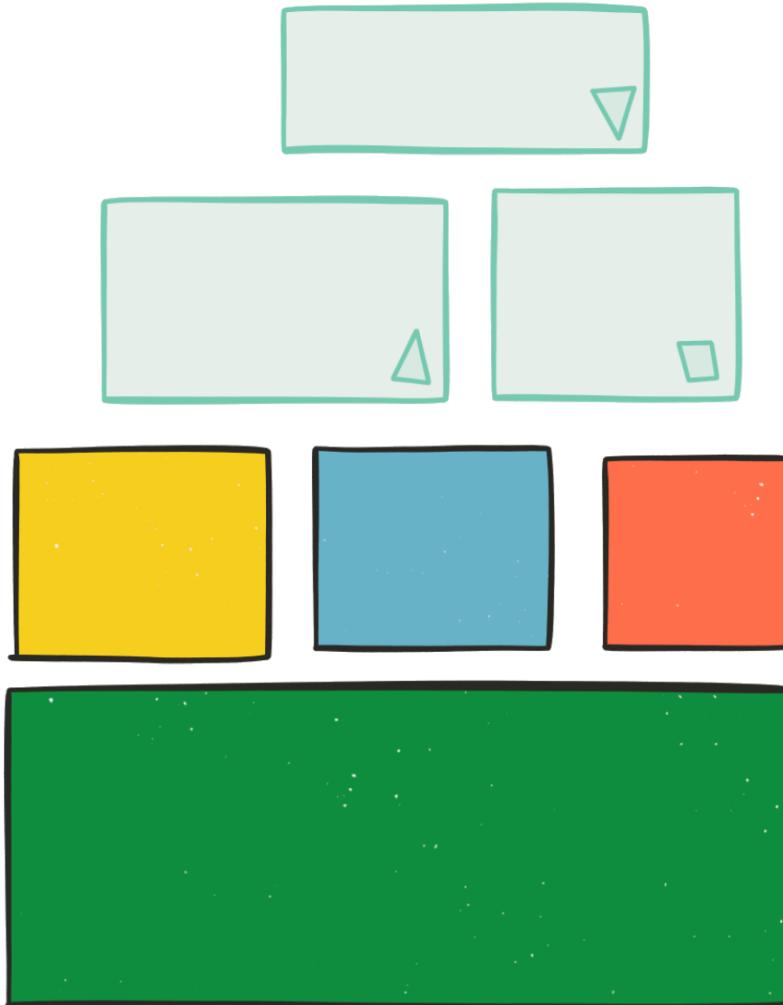
Lightweight



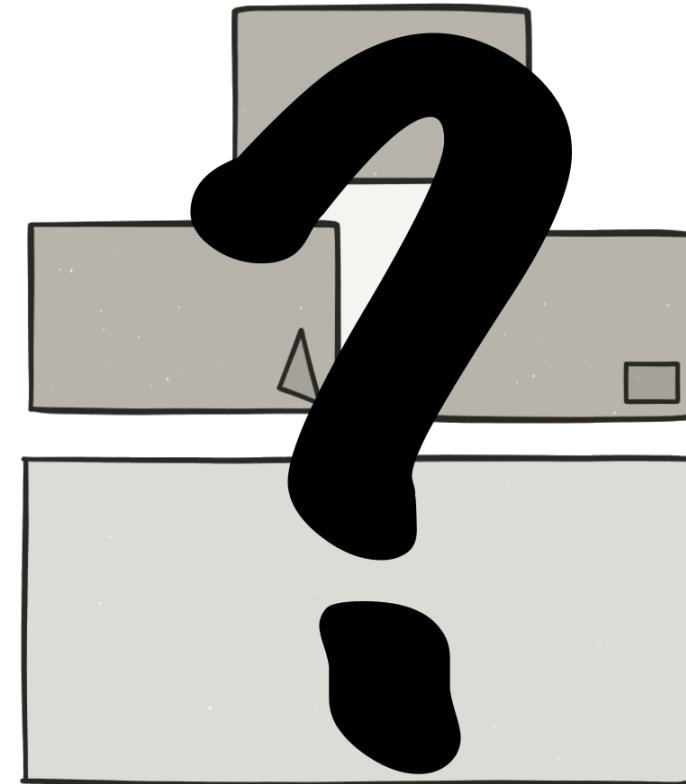
Composable



Getting serious ...



Multiple hosts deployments



Some challenges

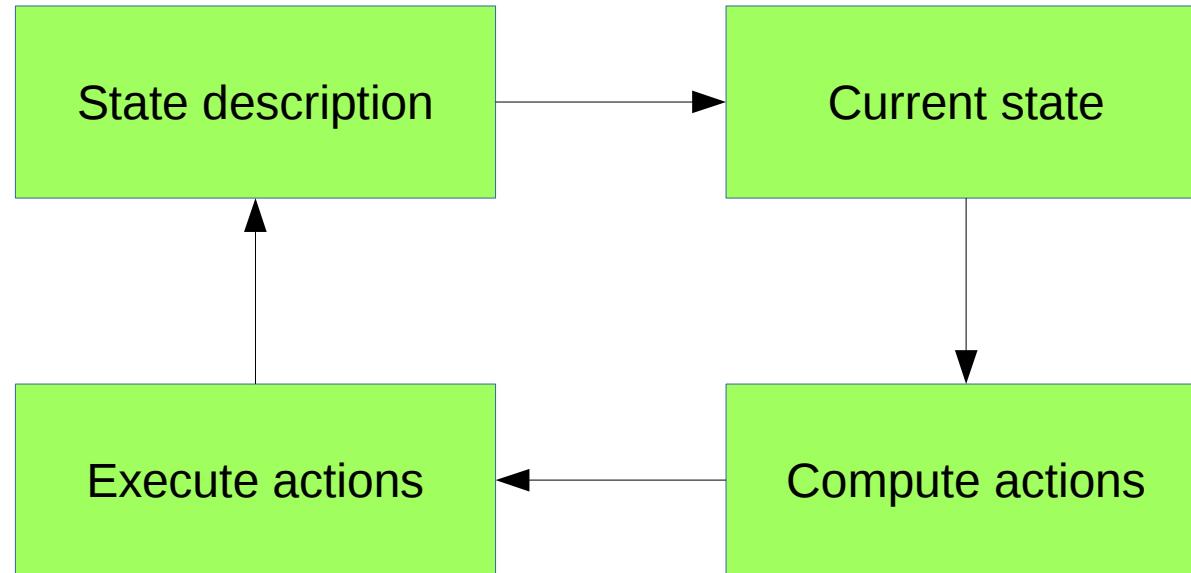
- Decide where to run each container
- Monitor nodes and containers
- Recover from failures
- Service discovery
- Expose services to external consumers
- Handle secrets (eg: credentials, certificates, keys,...)
- Handle data persistence

Containers orchestration

How can they help us?

- We don't have to find the right placement for each container
- We don't have to recover from failures manually
- We just declare a desired state

Desired state reconciliation



Which one?



Nomad



kubernetes



MESOS



Kubernetes, of course!!!

Kubernetes (aka k8s)

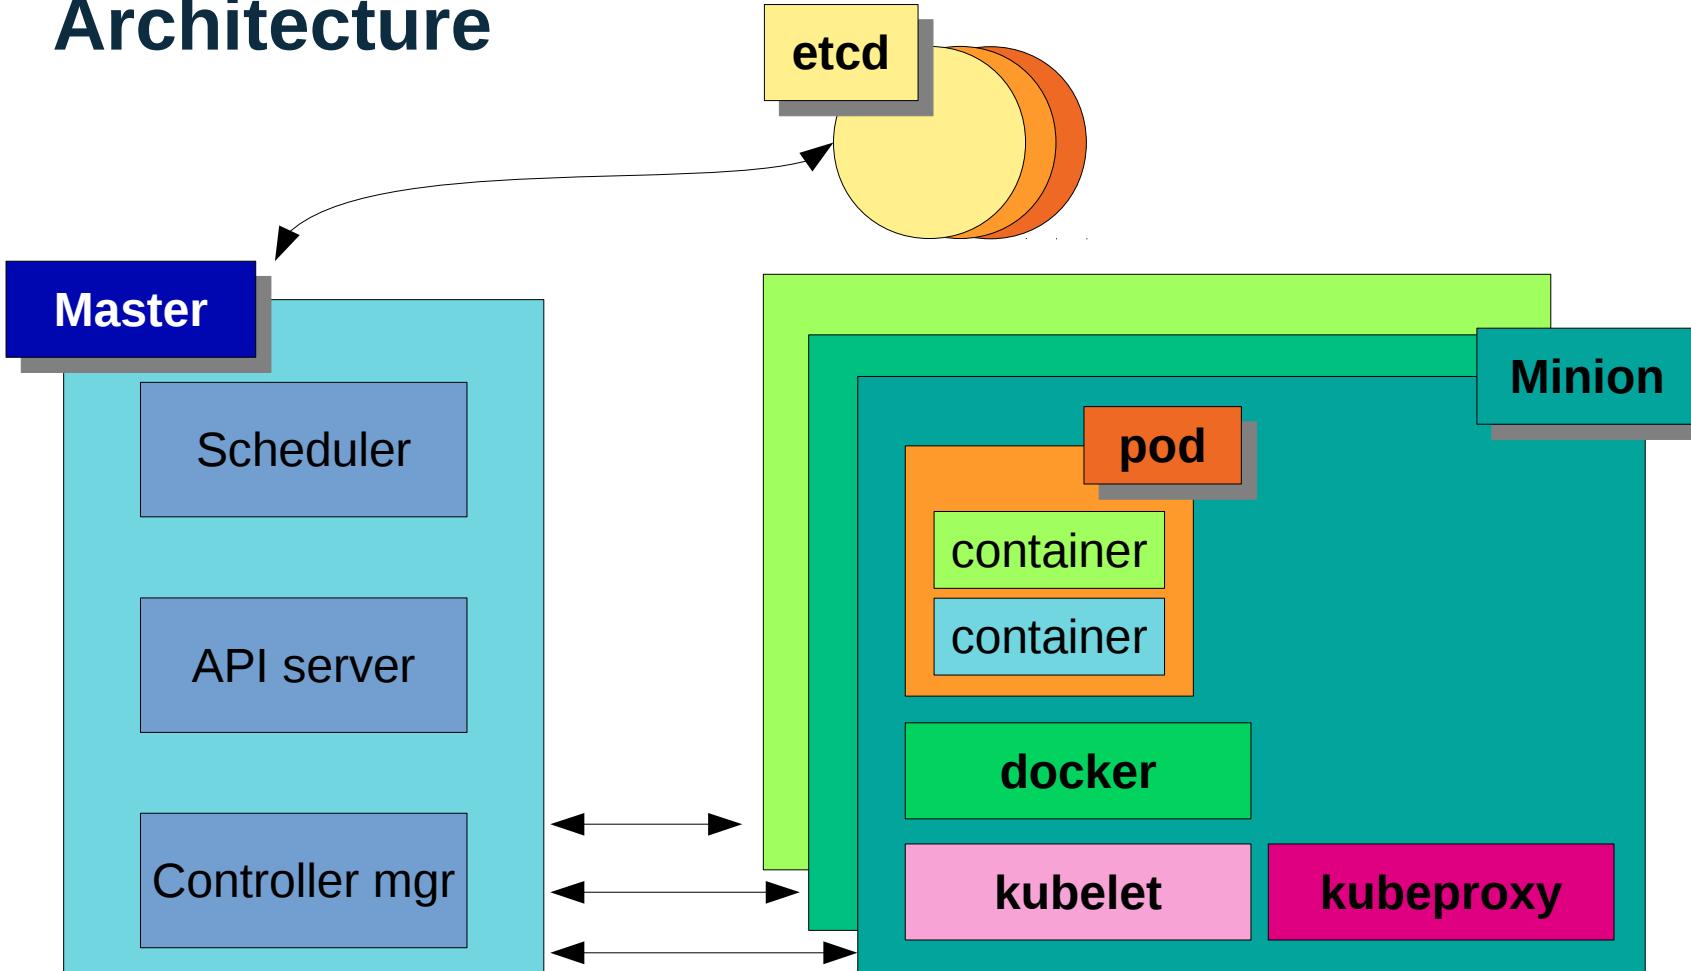
- Created by Google, now part of CNCF
- Solid design
- Big and active community
- An established “brand”
- Opinionated solution, has an answer to most questions



Kubernetes advantages

- Workload portability: doesn't enforce its directives to the application.
- Friendly with legacy applications: smooth migration path.
- Avoid vendor lock-in.
- Self healing.
- Auto-scaling.
- Has a solution for many problems:
 - Persistent storage.
 - Secrets management.
 - Blue-Green deployments.
- Flexible: plug-in architecture

Architecture



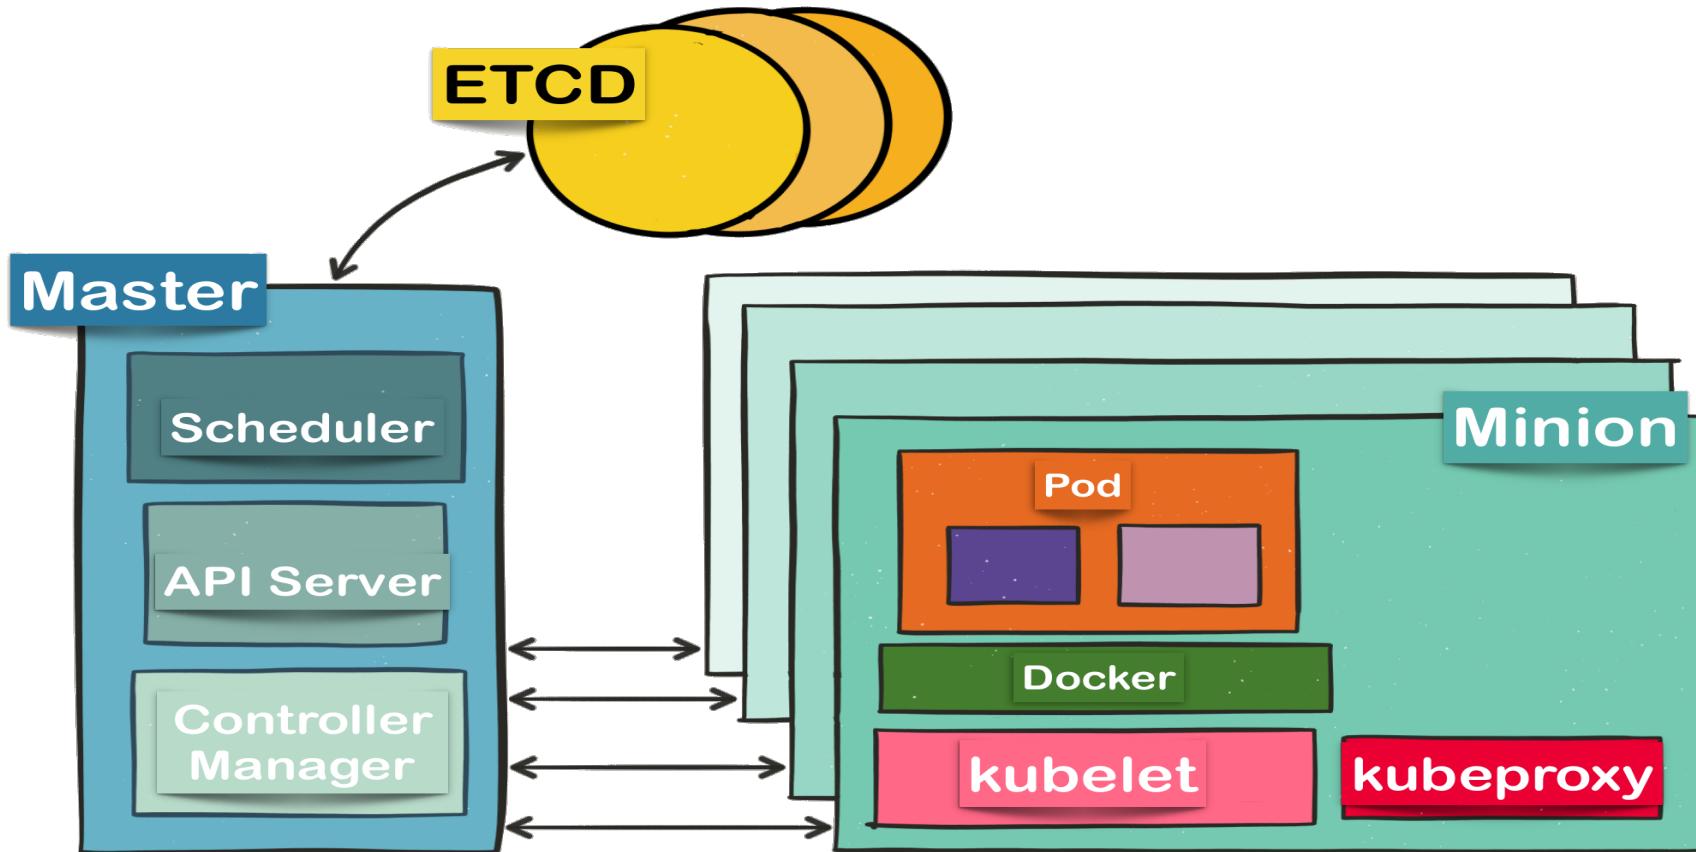
Kubernetes concepts

- Containers
- Pods
- Replication controllers
- Services
- Namespaces
- Labels
- Volumes

Kubernetes deployment

- Requires an *etcd* cluster.
- Requires one or more *master* nodes.
- Requires one or more *worker* nodes.
- Requires a SDN network joining all the *worker* nodes.
- Requires a load balancer to expose internal applications.
- Lots of patience to link all these components together.

Architecture



Kubernetes status

- It's a pleasure to use as developer deploying your application.
- It's a pleasure to administer as an operator.
- It's a pain to deploy.

Things are changing

- **kubeadm**: upstream tool for kubernetes deployment.
- Kubernetes deployed with containers, by kubernetes itself.
- Introduced with the 1.4 release, still alpha.
- External orchestrators
 - OpenStack Magnum
 - Salt
 - Terraform

Replication Controller

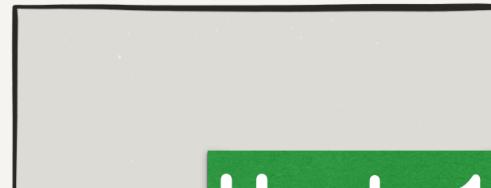
Application



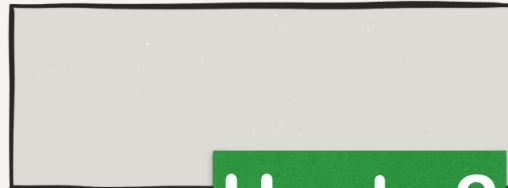
Web



DB



Host 1



Host 2

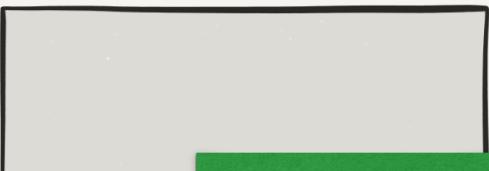
Application



Web



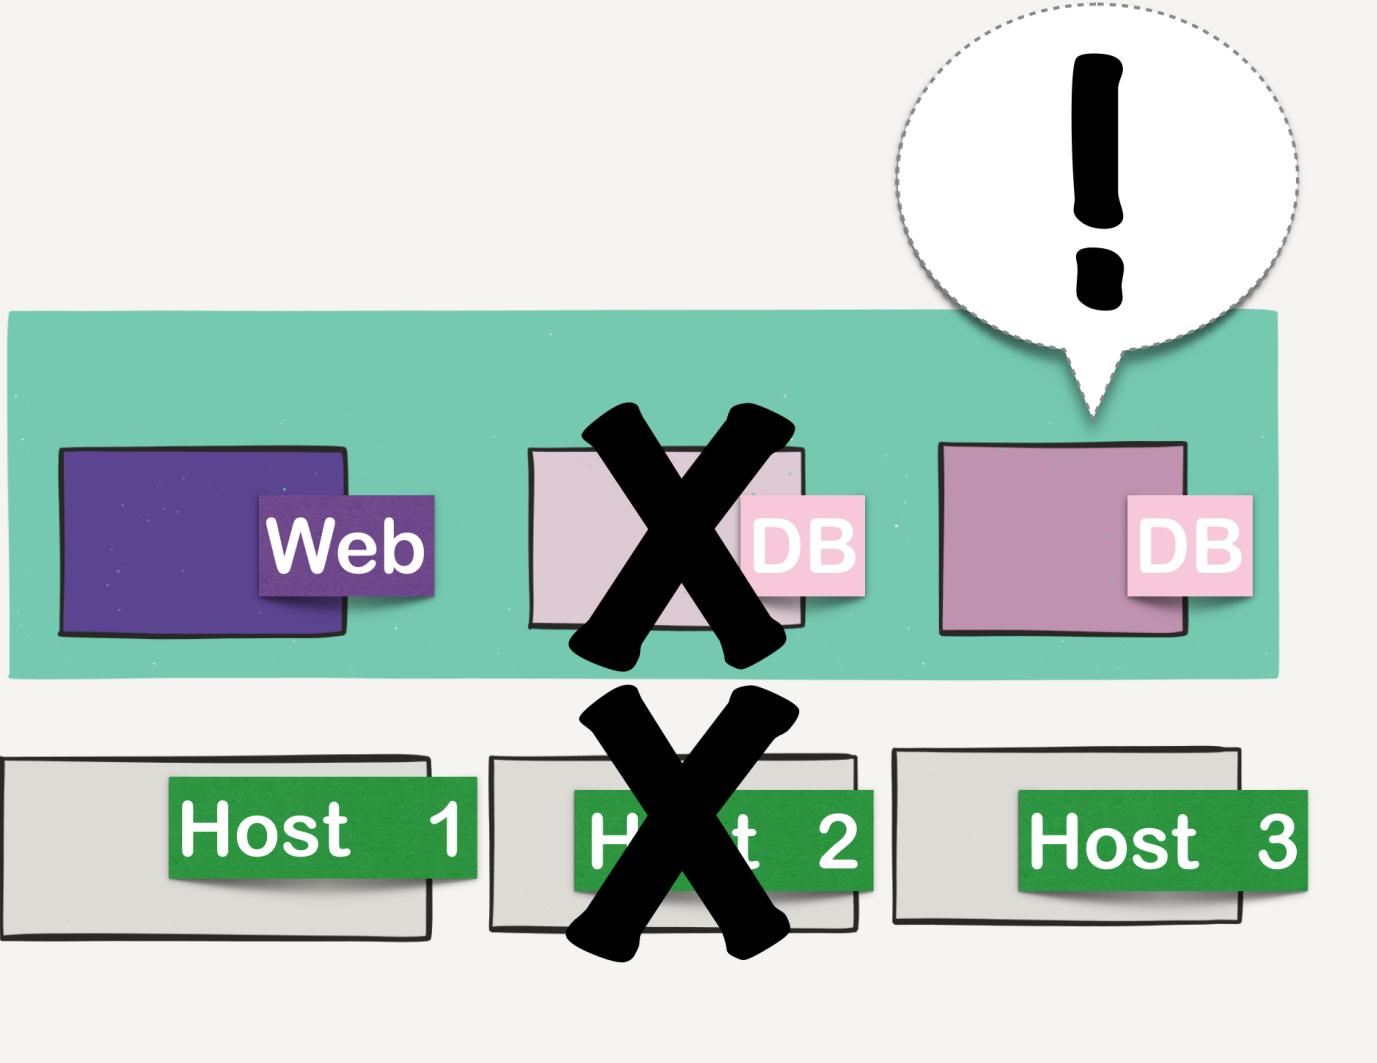
DB



Host 1

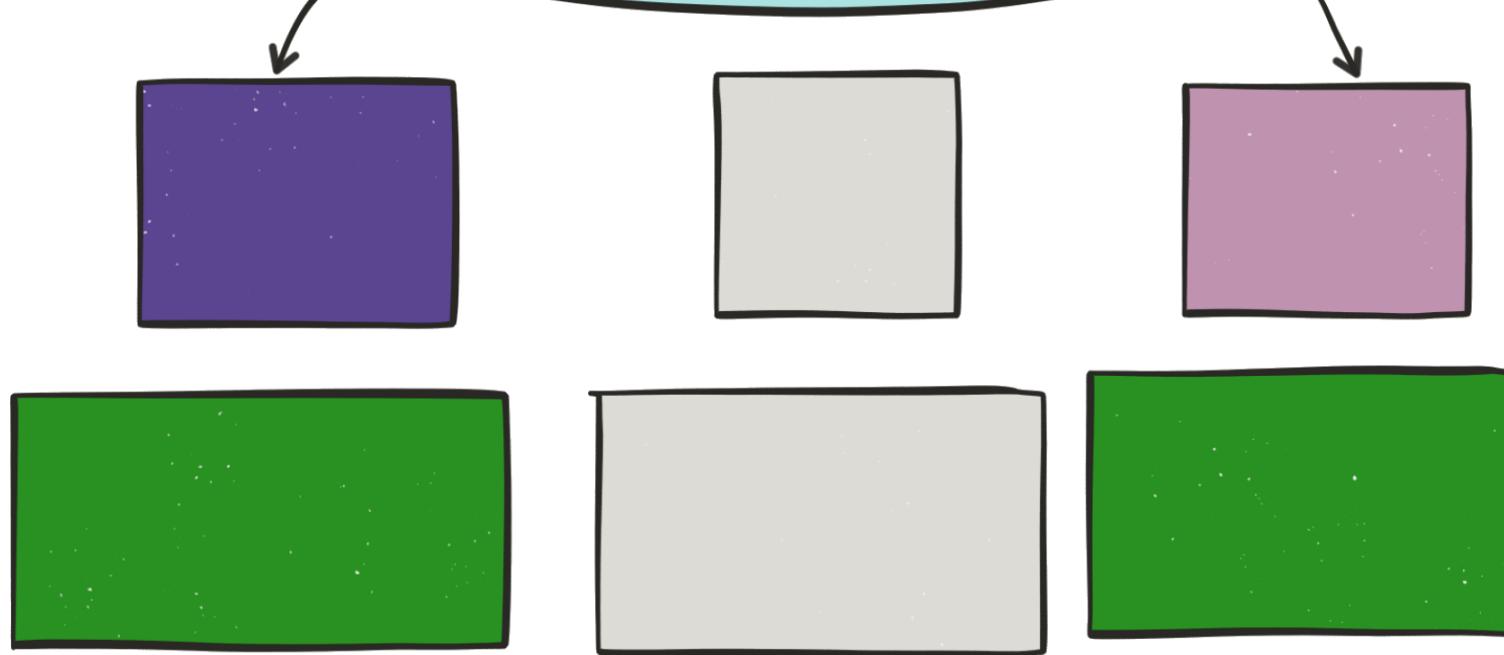


Host 2



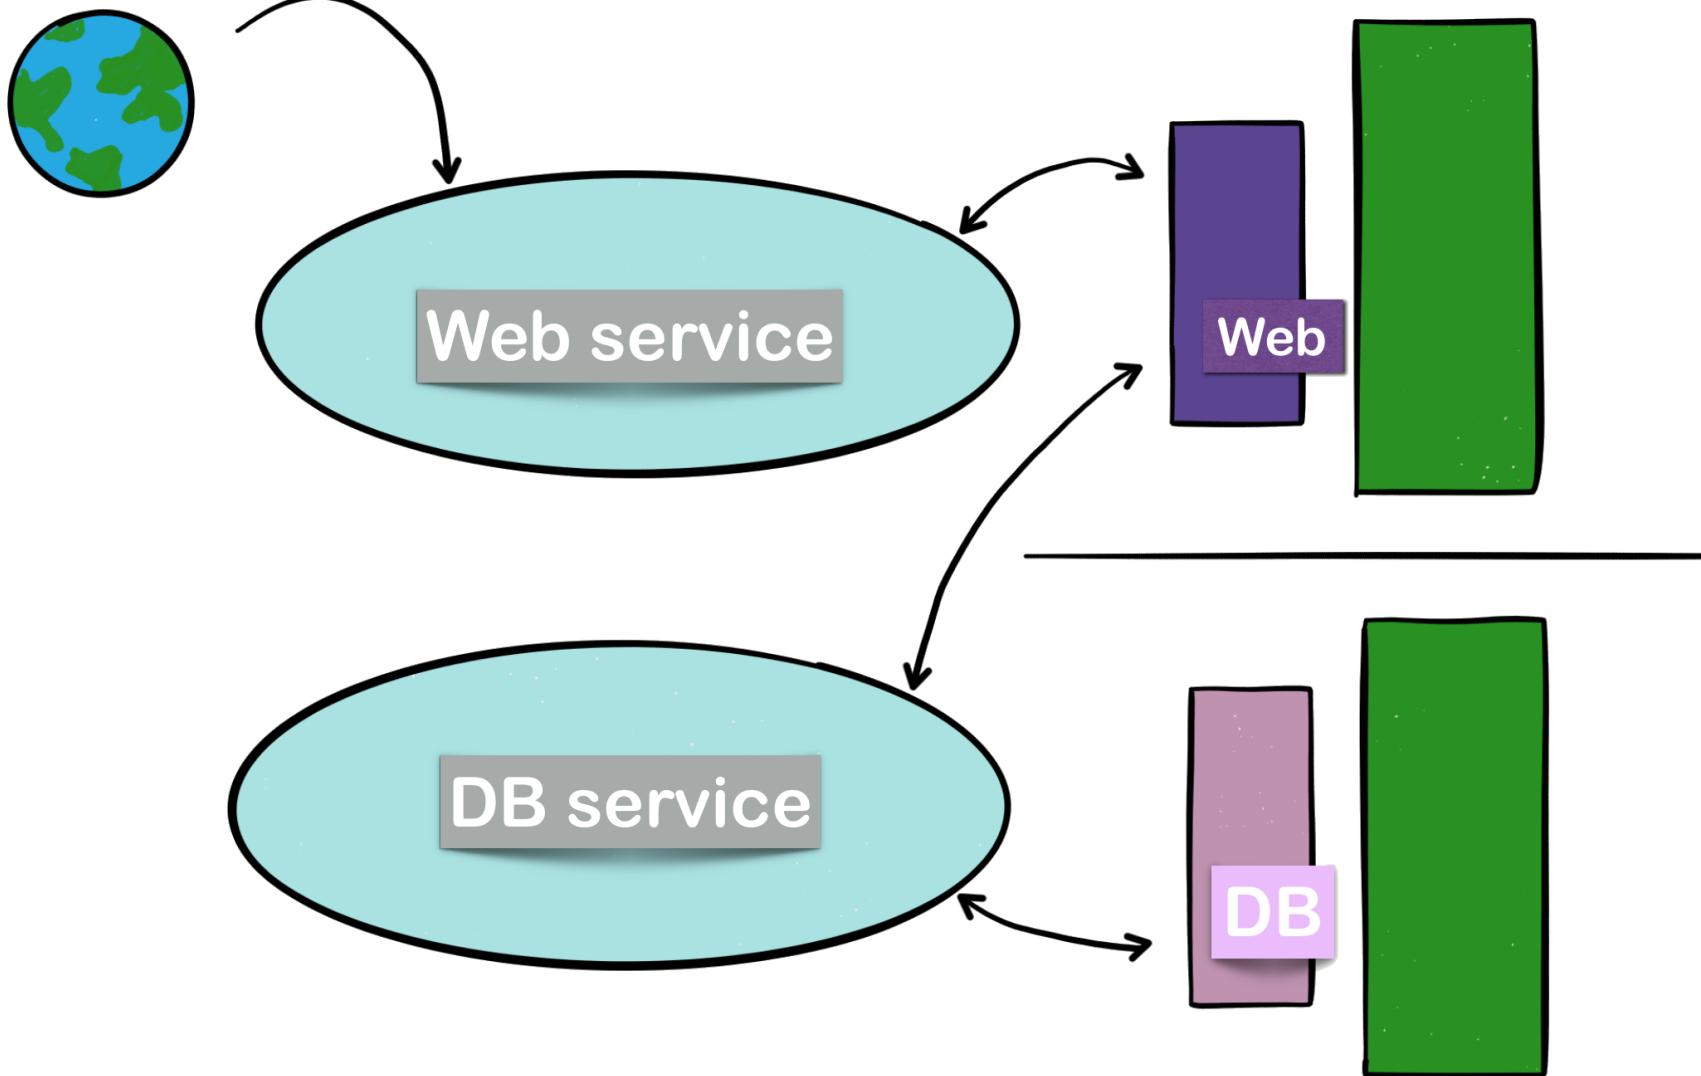
Kubernetes Services

Service



Basically

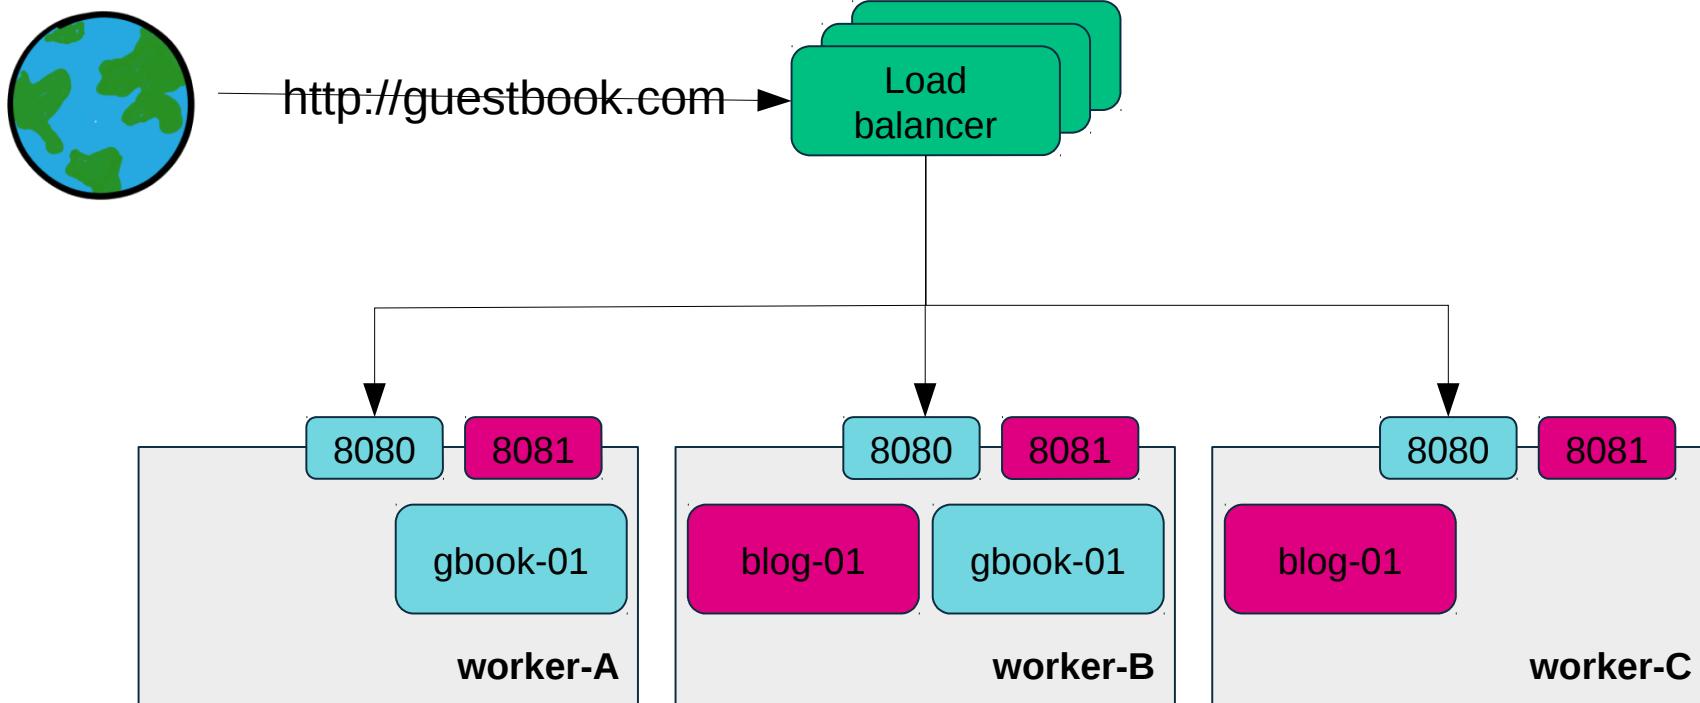




Implementation details

- Internal service:
 - Virtual IP: defined at creation time, doesn't change
 - Internal load balancer: kube-proxy
 - An internal DNS can point to the Virtual IP
- Published service: like an internal one, but...
 - Each minion exposes the internal service on a port
 - Can use cloud providers to have a load balancer

Architecture



What about data?

Data persistence

- Not all applications can offload data somewhere else
- Data should persist when the POD is moved to another host

Kubernetes Volumes

- Built into Kubernetes
- They are like other resources → scheduler awareness
- Support different backends via dedicated drivers:
 - Ceph
 - Cinder
 - NFS
 - GlusterFS
 -

Demo example



Questions ?

Michal Jura
Linux Cloud/HA Developer
mjura@suse.com