

---

---

# Reverse proxy automation - Varnish as a Service

Tomasz Ziółkowski 26.06.2018

---

# Overview



## Reverse Proxy

- Forward vs Reverse
- Features
- Caching

## Varnish

- Design
- Architecture
- State machine

## Allegro use case

- Some metrics
- Monolith to SOA
- ESI

---

# Overview

**allegro** Tech

## VaaS

- Overview
- Automation
- Contribute :)

---



---

# Reverse Proxy

---

---

# Proxy - reverse vs. forward



---

# Reverse proxy - features

## Hide architecture

Hide sophisticated architecture of microservices

## Access control on HTTP layer

L7 ACL based on http headers or client ip

## SSL termination

Terminate SSL connection before origin servers

---

---

# Reverse proxy - features

## Loadbalancing

Balance traffic based on several predefined LB algorithms

## Link usage optimization

Compress content for a client (if needed)

## ESI

Composite html responses for a client from several origin servers responses

---

---

# Reverse proxy - features

## Caching

Content caching based on Cache-Control response headers


## Availability

Active checks

## Scalability

Reconfigure proxy after autoscaling

Highly {  
Reliable  
Scalable  
Available





---

# Reverse proxy - caching

## Cache-Control response header

Caching levels

- public
- private
- no-cache
- no-store

**Cache-Control:** public,max-age=420;

---

---

# Reverse proxy - caching

## ETag response header

The ETag HTTP response header is an identifier for a specific version of a resource. It allows caches to be more efficient, and saves bandwidth, as a web server does not need to send a full response if the content has not changed.

**ETag:** "737060cd8c284d8af7ad3082f209582d"

## If-None-Match response header

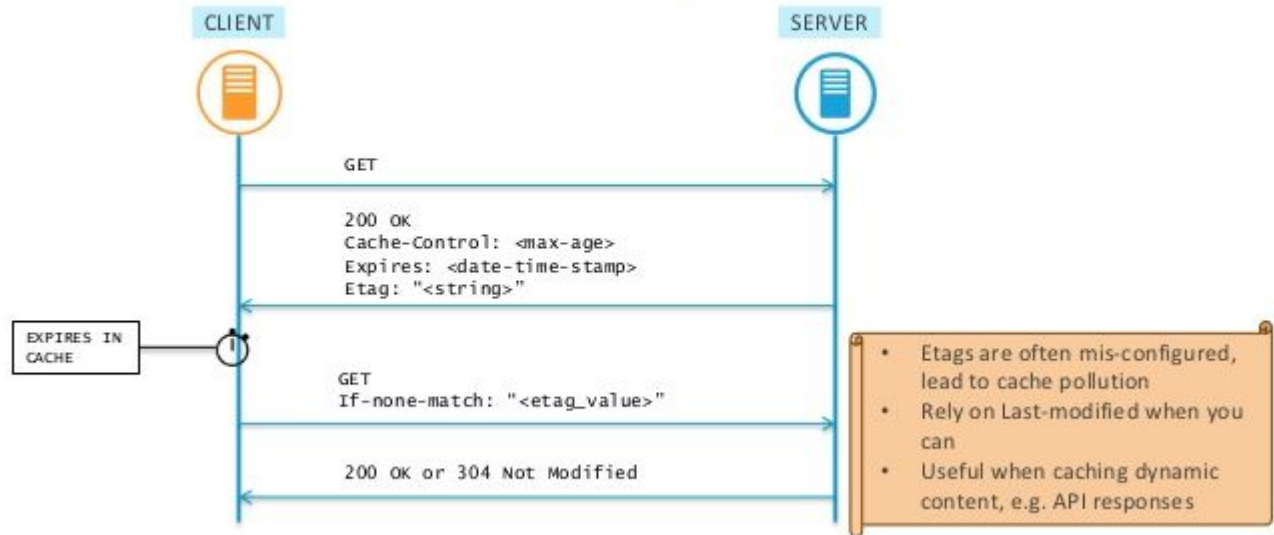
Check if outdated resource are really modified.

**If-None-Match:** "737060cd8c284d8af7ad3082f209582d"

---

# Reverse proxy - caching

## TYPICAL CACHE INTERACTION (ETAG)



---

# Reverse proxy - caching

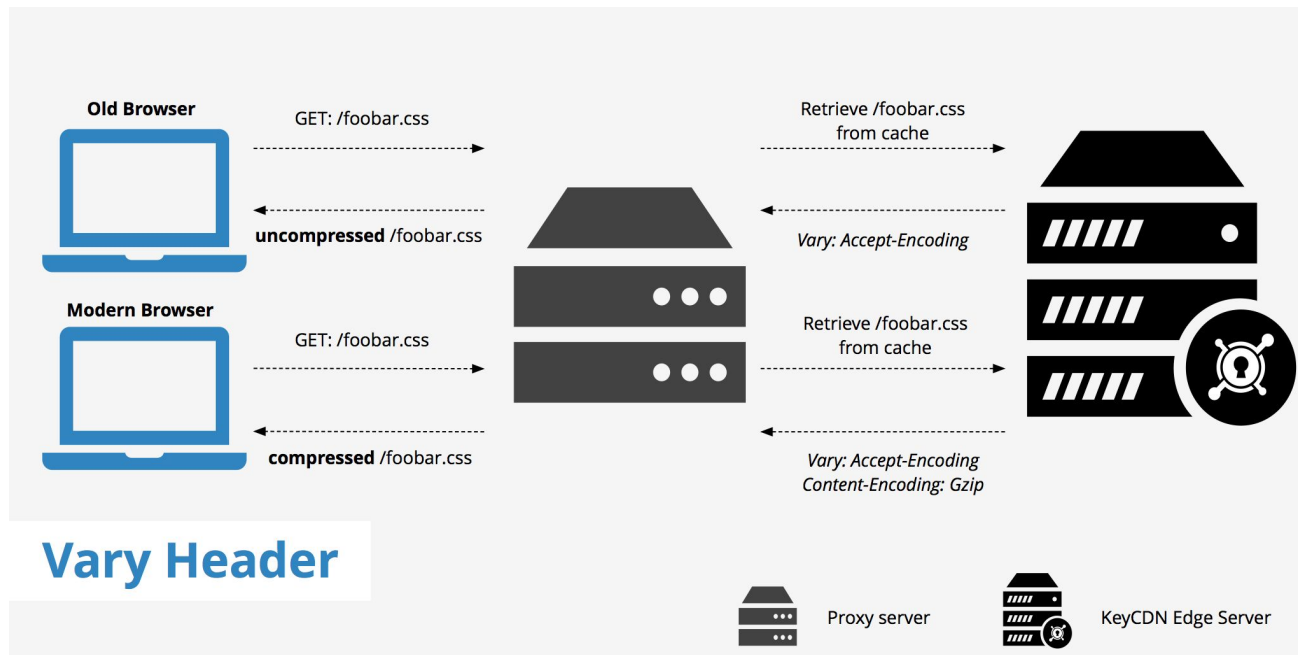
## Vary response header

The Vary HTTP response header determines how to match future request headers to decide whether a cached response can be used rather than requesting a fresh one from the origin server. It is used by the server to indicate which headers it used when selecting a representation of a resource in a content negotiation algorithm.

**Vary:** Accept-Encoding

---

# Reverse proxy - caching





# Varnish

---

---

# Varnish - design

Designed as reverse-proxy

Handle only HTTP traffic

Configurable via DSL

No native API (CLI via telnet)

---

---

# Varnish - terminology

Backend - ip:port - origin server

Director - logical set of backends enriched with LB algorithms (RR, Random, Hash)

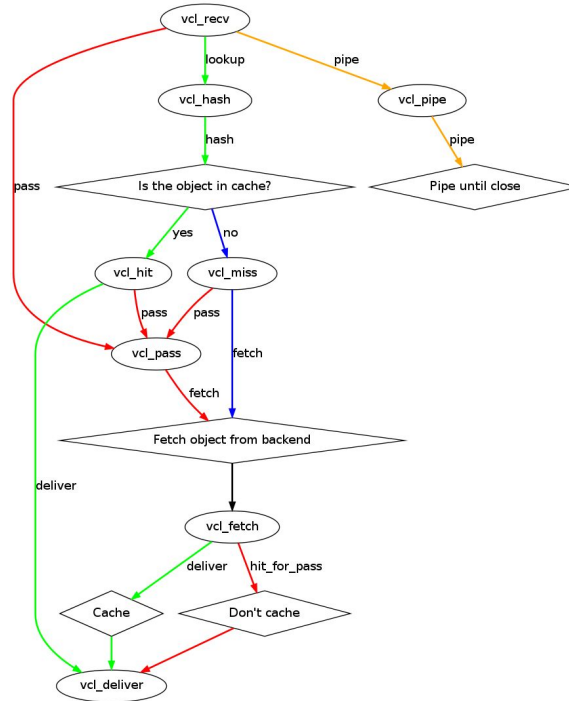
Probe - verify backend readiness to handle traffic

VCL - Varnish Configuration Language


---



# Varnish - state machine



# Varnish - simple flow



---

# Varnish

Possibility to expand state machine

Possibility to modify http headers req/resp  
bereq/beresp

Possibility to use simple logical conditions

---

---

# Varnish - VCL

Directors/backend definition

Predefined sub implementation:

- Recv
- Fetch
- Deliver

Custom sub implementation

---

---

# Varnish - VCL example

```
backend default_backend {  
    .host = "127.0.0.1";  
    .port = "8888";  
}  
  
backend allegro {  
    .host = "allegro.pl";  
    .port = "80";  
}  
  
sub vcl_recv {  
    if (req.http.host == "allegro.pl") {  
        set req.backend_hint = allegro;  
    } else {  
        set req.backend_hint = default_backend;  
    }  
}  
  
sub vcl_deliver {  
    set resp.http.X-Hit = "HIT " + obj.hits;  
}
```

---

---

# Allegro use case

---

---

# allegro.pl - some metrics

20k rps

5Gbps

>100 front microservices

>800 origin servers


**allegro** Tech

---

---

# Allegro - monolith to SOA

Old http architecture /  
monolith app





---

# Allegro - monolith to SOA

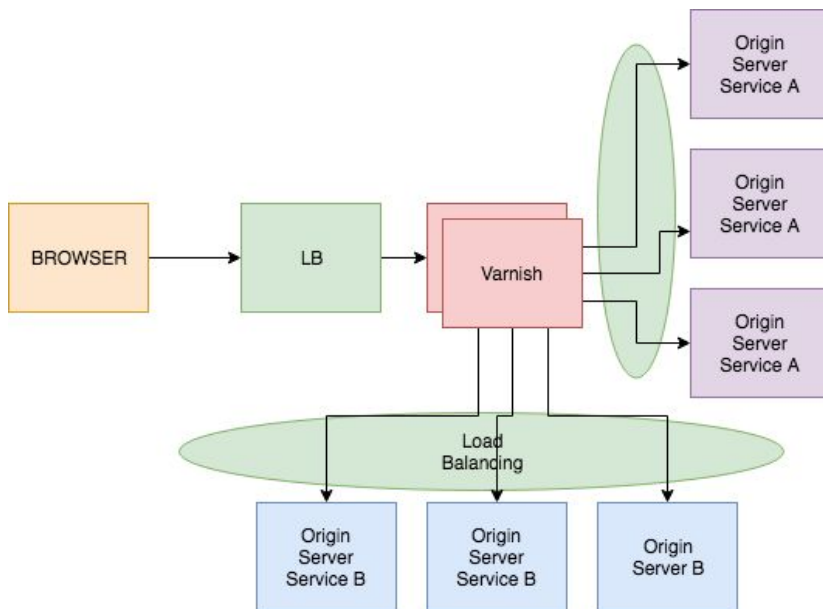
## Challenges:

- Application routing to multiple microservices
- Composite single html response based on several microservices
- Unhealthy instances - active detection
- Load balancing based on cookies
- Retries for 5xx response codes
- Move cache near to client

---

# Allegro - monolith to SOA

Http architecture - varnish /  
SOA



---


# Allegro - ESI

## Composite single html response:

- In main response use `<esi src="/path/to/fragment">`
- Varnish parse response and execute subrequest for each ESI tag
- ESI tags in main response are replaced by responses from subrequests

# Allegro - ESI

HOW DOES ESI WORK IN  
VARNISH-CACHE?



Slide 18 - Presenter view



---

# VaaS - motivation

Monolith to SOA

Lack of management tools


Reload configuration on each change

Automated reload after automated change

Automated support for local processing

---

# VaaS - overview



---

# VaaS - automation

Generating VCL based on templates & DB objects representing directors and backends

Homogeneous clusters (same configuration on each machine in the same dc)

Directors & backends exists independently of VCL

---



---

# VaaS - automation

Automated generating VCL regard to local processing

Synchronous reconfiguration

Expose REST API - enable automation with other ecosystem components (ie. service-discovery)

---

---

# VaaS - automation - Templates

## Placeholders:

- VCL
  - DIRECTORS
  - ROUTER
  - SET\_BACKEND\_director\_name
-

---

# VaaS - automation - Templates

```
<VCL>
  <HEADERS/>
  <ACL/>
  <DIRECTORS>
    <DIRECTOR_{DIRECTOR}>
      <BACKEND_DEFINITION_LIST_{DIRECTOR}_{DC}/>
      <DIRECTOR_DEFINITION__{DIRECTOR}_{DC}>
        <BACKEND_LIST_{DIRECTOR}_{DC}/>
      </DIRECTOR_DEFINITION__{DIRECTOR}_{DC}>
    </DIRECTOR_{DIRECTOR}_{DC}>
    ...
  </DIRECTORS>
  <RECV>
    <PROPER_PROTOCOL_REDIRECT/>
    <ROUTER>
      <SET_BACKEND_{DIRECTOR}/>
      ...
      <SET_BACKEND_{DIRECTOR}/>
    </ROUTER>
  </RECV>
  <OTHER_FUNCTIONS/>
</VCL>
```

---

---

---

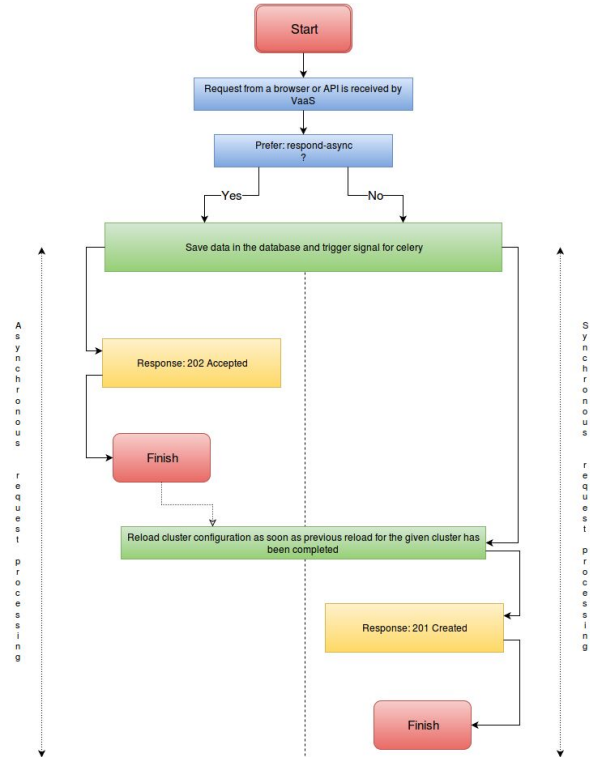
# VaaS - automation

VaaS is used by other parts of ecosystem via REST API

VaaS provide sync and async model of making changes

---

# VaaS - automation



---

# VaaS - automation

- Application routing - SOA
- Front microservices deployment
- Canary
- Redirect protocols
- Composite site from fragments
- HA, cross-dc fallbacks

... any many more is easier thanks to Varnish+VaaS

---

---

# VaaS numbers - in allegro

- 20 clusters
  - >70 varnishes
  - >200 directors
  - >1000 backends
  - 50-100 deployments per day
  - Thousands of VCL changes per day
-

---

# VaaS - allegro dev perspective

- Integrated with sophisticated ecosystem
  - Transparent for most operations
  - Similar to declarative infrastructure as LoadBalancer in K8s
-



---

# VaaS - contribute :)

## Prerequisite

VirtualBox

Vagrant

## Run development environment

git clone <https://github.com/allegro/vaas.git>

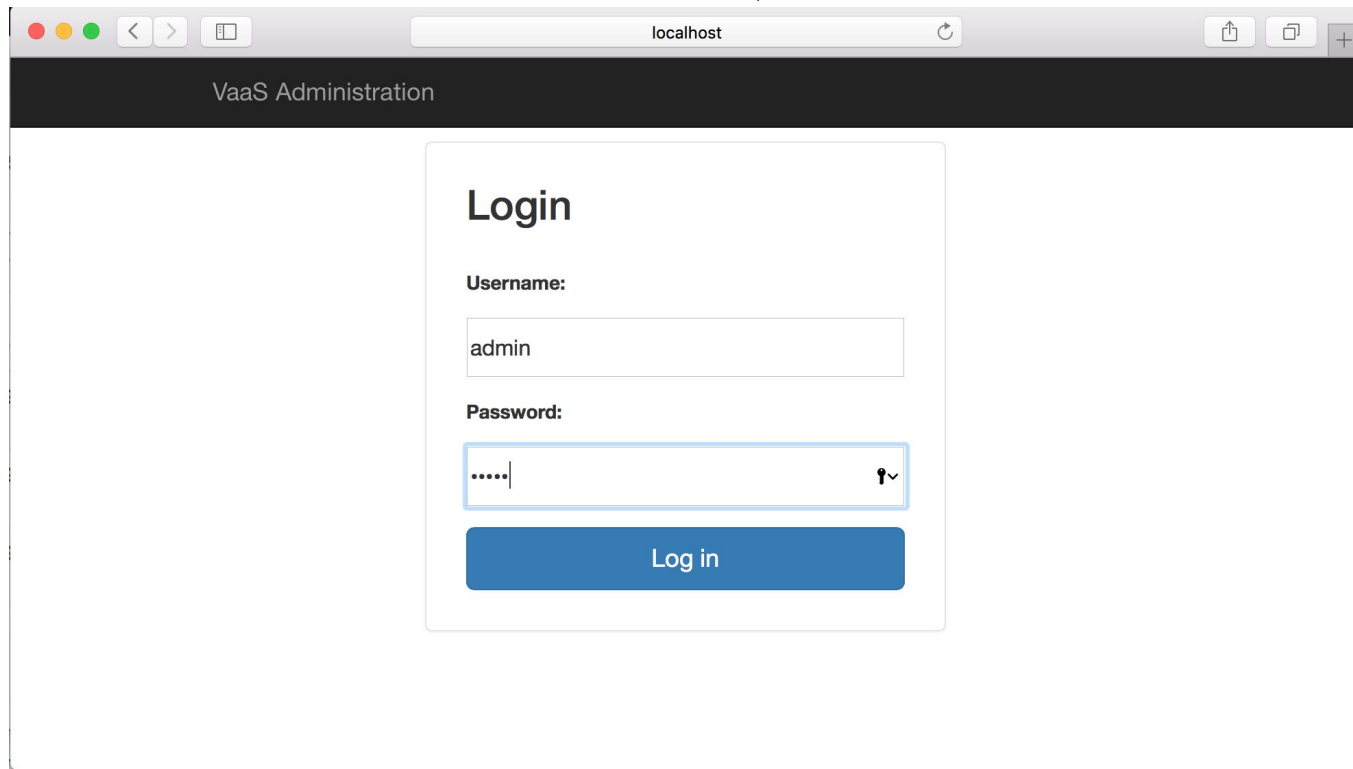
cd vaas

vagrant up

open <http://localhost:3030>

---

# VaaS - contribute :)



A screenshot of a web browser window displaying the VaaS Administration login interface. The browser's address bar shows 'localhost'. The page has a dark header with the text 'VaaS Administration'. The main content area is white and contains a centered login form. The form has a title 'Login', a 'Username:' label with a text input field containing 'admin', a 'Password:' label with a password input field showing masked characters '.....' and a toggle icon, and a blue 'Log in' button at the bottom.

VaaS Administration

## Login

Username:

Password:

Log in

# VaaS - contribute :)

The screenshot displays the VaaS Administration web interface. At the top, a browser window shows the address bar with 'localhost'. The page header includes 'VaaS Administration' on the left and 'Welcome, admin' and 'Recent Actions' on the right. A navigation bar below the header contains 'Site administration' and 'Applications'. The main content area is divided into two sections: 'Cluster' and 'Manager'. Each section contains a table with various configuration items, each having an 'Add' button and a 'Change' link.

VaaS Administration

Welcome, admin Recent Actions

Site administration Applications

## Cluster

Dcs	+ Add	Change
Logical clusters	+ Add	Change
Varnish servers	+ Add	Change
Vcl template blocks	+ Add	Change
Vcl templates	+ Add	Change
Vcl variables	+ Add	Change

## Manager

Backends	+ Add	Change
Directors	+ Add	Change
Probes	+ Add	Change
Time profiles	+ Add	Change
Purger		

# VaaS - contribute :)

VaaS Administration

Welcome, adminRecent Actions

Home / Cluster / Varnish servers

Select varnish server to changeAdd varnish server

SearchSearchFilter

Action: Go 0 of 3 selected

	Hostname	Ip	Port	Http port	Dc	Cluster	Cluster weight	Template	Template version	Enabled	Is connected	Canary	Vcl
-	varnish-4.1	192.168.199.4	6082	80	dc1	cluster2_siteB_test (2)	1	vagrant_template_4	4.0				Show vcl
-	varnish-4	192.168.199.3	6082	80	dc1	cluster2_siteB_test (2)	1	vagrant_template_4	4.0				Show vcl
-	varnish-3	192.168.199.2	6082	80	dc1	cluster1_siteA_test (1)	1	vagrant_template_3	3.0				Show vcl

---

# VaaS - useful links

1. <https://vaas.readthedocs.io/en/latest/>
  2. <https://github.com/allegro/vaas/>
  3. <https://allegro.tech/2015/09/vaas-open-source-management-tool-for-varnish-cache.html>
-

Q&A



---